

# Chapter 15

## Linux Data Structures



This appendix lists the major data structures that Linux uses and which are described in this book. They have been edited slightly to fit the paper.

### block\_dev\_struct

block\_dev\_struct data structures are used to register block devices as available for use by the buffer cache. They are held together in the blk\_dev vector.

```
struct blk_dev_struct {
    void (*request_fn)(void);
    struct request * current_request;
    struct request plug;
    struct tq_struct plug_tq;
};
```

### buffer\_head

The buffer\_head data structure holds information about a block buffer in the buffer cache.

```
/* bh state bits */
#define BH_Uptodate 0 /* 1 if the buffer contains valid data */
#define BH_Dirty 1 /* 1 if the buffer is dirty */
#define BH_Lock 2 /* 1 if the buffer is locked */
#define BH_Req 3 /* 0 if the buffer has been invalidated */
#define BH_Touched 4 /* 1 if the buffer has been touched (aging) */
#define BH_Has_aged 5 /* 1 if the buffer has been aged (aging) */
#define BH_Protected 6 /* 1 if the buffer is protected */
#define BH_FreeOnIO 7 /* 1 to discard the buffer_head after IO */

struct buffer_head {
    /* First cache line: */
    unsigned long b_blocknr; /* block number */
    kdev_t b_dev; /* device (B_FREE = free) */
    kdev_t b_rdev; /* Real device */
    unsigned long b_rsector; /* Real buffer location on disk */
    struct buffer_head *b_next; /* Hash queue list */
    struct buffer_head *b_this_page; /* circular list of buffers in one page */

    /* Second cache line: */
    unsigned long b_state; /* buffer state bitmap (above) */
    struct buffer_head *b_next_free;
    unsigned int b_count; /* users using this block */
    unsigned long b_size; /* block size */

    /* Non-performance-critical data follows. */
    char *b_data; /* pointer to data block */
    unsigned int b_list; /* List that this buffer appears */
    unsigned long b_flushtime; /* Time when this (dirty) buffer
                                * should be written */
    unsigned long b_lru_time; /* Time when this buffer was
```

```

                                * last used.                                */
struct wait_queue *b_wait;
struct buffer_head *b_prev;      /* doubly linked hash list      */
struct buffer_head *b_prev_free; /* doubly linked list of buffers */
struct buffer_head *b_reqnext;   /* request queue               */
};

```

## device

Every network device in the system is represented by a device data structure.

```

struct device
{
    /*
     * This is the first field of the "visible" part of this structure
     * (i.e. as seen by users in the "Space.c" file). It is the name
     * the interface.
     */
    char *name;

    /* I/O specific fields */
    unsigned long rmem_end;      /* shmem "recv" end */
    unsigned long rmem_start;    /* shmem "recv" start */
    unsigned long mem_end;       /* shared mem end */
    unsigned long mem_start;     /* shared mem start */
    unsigned long base_addr;     /* device I/O address */
    unsigned char irq;           /* device IRQ number */

    /* Low-level status flags. */
    volatile unsigned char start, /* start an operation */
                          interrupt; /* interrupt arrived */
    unsigned long tbusy;          /* transmitter busy */
    struct device *next;

    /* The device initialization function. Called only once. */
    int (*init)(struct device *dev);

    /* Some hardware also needs these fields, but they are not part of
       the usual set specified in Space.c. */
    unsigned char if_port;        /* Selectable AUI,TP, */
    unsigned char dma;            /* DMA channel */

    struct enet_statistics* (*get_stats)(struct device *dev);

    /*
     * This marks the end of the "visible" part of the structure. All
     * fields hereafter are internal to the system, and may change at
     * will (read: may be cleaned up at will).
     */

    /* These may be needed for future network-power-down code. */
    unsigned long trans_start;    /* Time (jiffies) of
                                   last transmit */
    unsigned long last_rx;        /* Time of last Rx */
    unsigned short flags;         /* interface flags (BSD) */
    unsigned short family;        /* address family ID */
    unsigned short metric;        /* routing metric */
    unsigned short mtu;           /* MTU value */
    unsigned short type;          /* hardware type */
    unsigned short hard_header_len; /* hardware hdr len */
    void *priv;                   /* private data */

    /* Interface address info. */
    unsigned char broadcast[MAX_ADDR_LEN];
    unsigned char pad;
    unsigned char dev_addr[MAX_ADDR_LEN];
    unsigned char addr_len;       /* hardware addr len */
}

```

```

unsigned long      pa_addr;          /* protocol address      */
unsigned long      pa_brdaddr;       /* protocol broadcast addr*/
unsigned long      pa_dstaddr;       /* protocol P-P other addr*/
unsigned long      pa_mask;          /* protocol netmask      */
unsigned short     pa_alen;          /* protocol address len  */

struct dev_mc_list *mc_list;         /* M'cast mac addrs      */
int               mc_count;          /* No installed mcasts   */

struct ip_mc_list  *ip_mc_list;      /* IP m'cast filter chain */
__u32              tx_queue_len;     /* Max frames per queue  */

/* For load balancing driver pair support */
unsigned long      pkt_queue;        /* Packets queued        */
struct device      *slave;           /* Slave device           */
struct net_alias_info *alias_info;   /* main dev alias info    */
struct net_alias   *my_alias;        /* alias devs             */

/* Pointer to the interface buffers. */
struct sk_buff_head buffs[DEV_NUMBUFFS];

/* Pointers to interface service routines. */
int (*open)(struct device *dev);
int (*stop)(struct device *dev);
int (*hard_start_xmit)(struct sk_buff *skb,
                      struct device *dev);
int (*hard_header)(struct sk_buff *skb,
                  struct device *dev,
                  unsigned short type,
                  void *daddr,
                  void *saddr,
                  unsigned len);
int (*rebuild_header)(void *eth,
                    struct device *dev,
                    unsigned long raddr,
                    struct sk_buff *skb);
void (*set_multicast_list)(struct device *dev);
int (*set_mac_address)(struct device *dev,
                      void *addr);
int (*do_ioctl)(struct device *dev,
                struct ifreq *ifr,
                int cmd);
int (*set_config)(struct device *dev,
                  struct ifmap *map);
void (*header_cache_bind)(struct hh_cache **hhp,
                        struct device *dev,
                        unsigned short htype,
                        __u32 daddr);
void (*header_cache_update)(struct hh_cache *hh,
                          struct device *dev,
                          unsigned char * haddr);
int (*change_mtu)(struct device *dev,
                 int new_mtu);
struct iw_statistics* (*get_wireless_stats)(struct device *dev);
};

```

## device\_struct

device\_struct data structures are used to register character and block devices (they hold its name and the set of file operations that can be used for this device). Each valid member of the chrdevs and blkdevs vectors represents a character or block device respectively.

```

struct device_struct {
    const char * name;
    struct file_operations * fops;
};

```

## file

Each open file, socket etcetera is represented by a file data structure.

```
struct file {
    mode_t f_mode;
    loff_t f_pos;
    unsigned short f_flags;
    unsigned short f_count;
    unsigned long f_reada, f_ramax, f_raend, f_ralen, f_rawin;
    struct file *f_next, *f_prev;
    int f_owner;          /* pid or -pgrp where SIGIO should be sent */
    struct inode * f_inode;
    struct file_operations * f_op;
    unsigned long f_version;
    void *private_data; /* needed for tty driver, and maybe others */
};
```

## files\_struct

The files\_struct data structure describes the files that a process has open.

```
struct files_struct {
    int count;
    fd_set close_on_exec;
    fd_set open_fds;
    struct file * fd[NR_OPEN];
};
```

## fs\_struct

```
struct fs_struct {
    int count;
    unsigned short umask;
    struct inode * root, * pwd;
};
```

## gendisk

The gendisk data structure holds information about a hard disk. They are used during initialization when the disks are found and then probed for partitions.

```
struct hd_struct {
    long start_sect;
    long nr_sects;
};

struct gendisk {
    int major;          /* major number of driver */
    const char *major_name; /* name of major driver */
    int minor_shift;    /* number of times minor is shifted to
                        get real minor */
    int max_p;          /* maximum partitions per device */
    int max_nr;         /* maximum number of real devices */

    void (*init)(struct gendisk *);
    /* Initialization called before we
    do our thing */
    struct hd_struct *part; /* partition table */
    int *sizes;             /* device size in blocks, copied to
                        blk_size[] */
    int nr_real;            /* number of real devices */
    void *real_devices;    /* internal use */
};
```

```
    struct gendisk *next;
};
```

## inode

The VFS inode data structure holds information about a file or directory on disk.

```
struct inode {
    kdev_t                i_dev;
    unsigned long         i_ino;
    umode_t               i_mode;
    nlink_t               i_nlink;
    uid_t                 i_uid;
    gid_t                 i_gid;
    kdev_t                i_rdev;
    off_t                 i_size;
    time_t                i_atime;
    time_t                i_mtime;
    time_t                i_ctime;
    unsigned long         i_blksize;
    unsigned long         i_blocks;
    unsigned long         i_version;
    unsigned long         i_nrpages;
    struct semaphore      i_sem;
    struct inode_operations *i_op;
    struct super_block    *i_sb;
    struct wait_queue     *i_wait;
    struct file_lock      *i_flock;
    struct vm_area_struct *i_mmap;
    struct page           *i_pages;
    struct dquot          *i_dquot[MAXQUOTAS];
    struct inode          *i_next, *i_prev;
    struct inode          *i_hash_next, *i_hash_prev;
    struct inode          *i_bound_to, *i_bound_by;
    struct inode          *i_mount;
    unsigned short        i_count;
    unsigned short        i_flags;
    unsigned char         i_lock;
    unsigned char         i_dirt;
    unsigned char         i_pipe;
    unsigned char         i_sock;
    unsigned char         i_seek;
    unsigned char         i_update;
    unsigned short        i_writecount;
    union {
        struct pipe_inode_info pipe_i;
        struct minix_inode_info minix_i;
        struct ext_inode_info ext_i;
        struct ext2_inode_info ext2_i;
        struct hpfs_inode_info hpfs_i;
        struct msdos_inode_info msdos_i;
        struct umsdos_inode_info umsdos_i;
        struct iso_inode_info isofs_i;
        struct nfs_inode_info nfs_i;
        struct xiafs_inode_info xiafs_i;
        struct sysv_inode_info sysv_i;
        struct affs_inode_info affs_i;
        struct ufs_inode_info ufs_i;
        struct socket         socket_i;
        void                  *generic_ip;
    } u;
};
```

## ipc\_perm

The ipc\_perm data structure describes the access permissions of a System V IPC object .

```
struct ipc_perm
{
    key_t    key;
    ushort  uid;    /* owner euid and egid */
    ushort  gid;
    ushort  cuid;   /* creator euid and egid */
    ushort  cgid;
    ushort  mode;   /* access modes see mode flags below */
    ushort  seq;   /* sequence number */
};
```

## irqaction

The irqaction data structure is used to describe the system's interrupt handlers.

```
struct irqaction {
    void (*handler)(int, void *, struct pt_regs *);
    unsigned long flags;
    unsigned long mask;
    const char *name;
    void *dev_id;
    struct irqaction *next;
};
```

## linux\_binfmt

Each binary file format that Linux understands is represented by a linux\_binfmt data structure.

```
struct linux_binfmt {
    struct linux_binfmt * next;
    long *use_count;
    int (*load_binary)(struct linux_binprm *, struct pt_regs * regs);
    int (*load_shlib)(int fd);
    int (*core_dump)(long signr, struct pt_regs * regs);
};
```

## mem\_map\_t

The mem\_map\_t data structure (also known as page) is used to hold information about each page of physical memory.

```
typedef struct page {
    /* these must be first (free area handling) */
    struct page *next;
    struct page *prev;
    struct inode *inode;
    unsigned long offset;
    struct page *next_hash;
    atomic_t count;
    unsigned flags;    /* atomic flags, some possibly
                        updated asynchronously */

    unsigned dirty:16,
               age:8;
    struct wait_queue *wait;
    struct page *prev_hash;
    struct buffer_head *buffers;
    unsigned long swap_unlock_entry;
    unsigned long map_nr;    /* page->map_nr == page - mem_map */
} mem_map_t;
```

## mm\_struct

The mm\_struct data structure is used to describe the virtual memory of a task or process.

```

struct mm_struct {
    int count;
    pgd_t * pgd;
    unsigned long context;
    unsigned long start_code, end_code, start_data, end_data;
    unsigned long start_brk, brk, start_stack, start_mmap;
    unsigned long arg_start, arg_end, env_start, env_end;
    unsigned long rss, total_vm, locked_vm;
    unsigned long def_flags;
    struct vm_area_struct * mmap;
    struct vm_area_struct * mmap_avl;
    struct semaphore mmap_sem;
};

```

## pci\_bus

Every PCI bus in the system is represented by a `pci_bus` data structure.

```

struct pci_bus {
    struct pci_bus *parent;      /* parent bus this bridge is on */
    struct pci_bus *children;    /* chain of P2P bridges on this bus */
    struct pci_bus *next;        /* chain of all PCI buses */

    struct pci_dev *self;        /* bridge device as seen by parent */
    struct pci_dev *devices;     /* devices behind this bridge */

    void *sysdata;              /* hook for sys-specific extension */

    unsigned char number;        /* bus number */
    unsigned char primary;       /* number of primary bridge */
    unsigned char secondary;     /* number of secondary bridge */
    unsigned char subordinate;   /* max number of subordinate buses */
};

```

## pci\_dev

Every PCI device in the system, including PCI-PCI and PCI-ISA bridge devices is represented by a `pci_dev` data structure.

```

/*
 * There is one pci_dev structure for each slot-number/function-number
 * combination:
 */
struct pci_dev {
    struct pci_bus *bus;         /* bus this device is on */
    struct pci_dev *sibling;     /* next device on this bus */
    struct pci_dev *next;        /* chain of all devices */

    void *sysdata;              /* hook for sys-specific extension */

    unsigned int devfn;          /* encoded device & function index */
    unsigned short vendor;
    unsigned short device;
    unsigned int class;          /* 3 bytes: (base,sub,prog-if) */
    unsigned int master : 1;     /* set if device is master capable */
    /*
     * In theory, the irq level can be read from configuration
     * space and all would be fine. However, old PCI chips don't
     * support these registers and return 0 instead. For example,
     * the Vision864-P rev 0 chip can use INTA, but returns 0 in
     * the interrupt line and pin registers. pci_init()
     * initializes this field with the value at PCI_INTERRUPT_LINE
     * and it is the job of pcibios_fixup() to change it if
     * necessary. The field must not be 0 unless the device
     * cannot generate interrupts at all.
     */
};

```

```
    unsigned char  irq;          /* irq generated by this device */
};
```

## request

request data structures are used to make requests to the block devices in the system. The requests are always to read or write blocks of data to or from the buffer cache.

```
struct request {
    volatile int  rq_status;
#define RQ_INACTIVE    (-1)
#define RQ_ACTIVE      1
#define RQ SCSI_BUSY    0xffff
#define RQ SCSI_DONE    0xfffe
#define RQ SCSI_DISCONNECTING  0xffe0

    kdev_t  rq_dev;
    int cmd;      /* READ or WRITE */
    int errors;
    unsigned long sector;
    unsigned long nr_sectors;
    unsigned long current_nr_sectors;
    char * buffer;
    struct semaphore * sem;
    struct buffer_head * bh;
    struct buffer_head * bhtail;
    struct request * next;
};
```

## rtable

Each rtable data structure holds information about the route to take in order to send packets to an IP host. rtable data structures are used within the IP route cache.

```
struct rtable
{
    struct rtable    *rt_next;
    __u32            rt_dst;
    __u32            rt_src;
    __u32            rt_gateway;
    atomic_t         rt_refcnt;
    atomic_t         rt_use;
    unsigned long    rt_window;
    atomic_t         rt_lastuse;
    struct hh_cache   *rt_hh;
    struct device     *rt_dev;
    unsigned short    rt_flags;
    unsigned short    rt_mtu;
    unsigned short    rt_irtt;
    unsigned char     rt_tos;
};
```

## semaphore

Semaphores are used to protect critical data structures and regions of code. y

```
struct semaphore {
    int count;
    int waking;
    int lock ;          /* to make waking testing atomic */
    struct wait_queue *wait;
};
```

## sk\_buff



The `sk_buff` data structure is used to describe network data as it moves between the layers of protocol.

```
struct sk_buff
{
    struct sk_buff      *next;          /* Next buffer in list          */
    struct sk_buff      *prev;          /* Previous buffer in list      */
    struct sk_buff_head *list;          /* List we are on               */
    int                  magic_debug_cookie;
    struct sk_buff      *link3;         /* Link for IP protocol level buffer chains */
    struct sock          *sk;           /* Socket we are owned by       */
    unsigned long        when;          /* used to compute rtt's        */
    struct timeval       stamp;         /* Time we arrived              */
    struct device        *dev;          /* Device we arrived on/are leaving by */
    union
    {
        struct tcphdr   *th;
        struct ethhdr   *eth;
        struct iphdr     *iph;
        struct udphdr   *uh;
        unsigned char    *raw;
        /* for passing file handles in a unix domain socket */
        void             *filp;
    } h;

    union
    {
        /* As yet incomplete physical layer views */
        unsigned char    *raw;
        struct ethhdr    *ethernet;
    } mac;

    struct iphdr         *ip_hdr;       /* For IPPROTO_RAW              */
    unsigned long         len;           /* Length of actual data        */
    unsigned long         csum;          /* Checksum                     */
    __u32                 saddr;        /* IP source address            */
    __u32                 daddr;        /* IP target address            */
    __u32                 raddr;        /* IP next hop address          */
    __u32                 seq;          /* TCP sequence number          */
    __u32                 end_seq;      /* seq [+ fin] [+ syn] + datalen */
    __u32                 ack_seq;      /* TCP ack sequence number      */
    unsigned char         proto_priv[16];
    volatile char         acked,        /* Are we acked ?               */
                        used,          /* Are we in use ?              */
                        free,          /* How to free this buffer      */
                        arp;           /* Has IP/ARP resolution finished */
    unsigned char         tries,        /* Times tried                  */
                        lock,          /* Are we locked ?              */
                        localroute,    /* Local routing asserted for this frame */
                        pkt_type,      /* Packet class                  */
                        pkt_bridged,   /* Tracker for bridging          */
                        ip_summed;     /* Driver fed us an IP checksum */

#define PACKET_HOST      0             /* To us                         */
#define PACKET_BROADCAST 1             /* To all                        */
#define PACKET_MULTICAST 2             /* To group                      */
#define PACKET_OTHERHOST 3            /* To someone else              */
    unsigned short        users;        /* User count - see datagram.c,tcp.c */
    unsigned short        protocol;     /* Packet protocol from driver.      */
    unsigned int          truesize;     /* Buffer size                      */
    atomic_t              count;        /* reference count                  */
    struct sk_buff        *data_skb;    /* Link to the actual data skb      */
    unsigned char         *head;        /* Head of buffer                  */
    unsigned char         *data;        /* Data head pointer               */
    unsigned char         *tail;        /* Tail pointer                    */
    unsigned char         *end;         /* End pointer                     */
    void                  (*destructor)(struct sk_buff *); /* Destruct function */
    __u16                 redirport;    /* Redirect port                   */
};
```

## sock

Each sock data structure holds protocol specific information about a BSD socket. For example, for an INET (Internet Address Domain) socket this data structure would hold all of the TCP/IP and UDP/IP specific information.

```
struct sock
{
    /* This must be first. */
    struct sock      *sklist_next;
    struct sock      *sklist_prev;

    struct options    *opt;
    atomic_t          wmem_alloc;
    atomic_t          rmem_alloc;
    unsigned long     allocation;      /* Allocation mode */
    __u32             write_seq;
    __u32             sent_seq;
    __u32             acked_seq;
    __u32             copied_seq;
    __u32             rcv_ack_seq;
    unsigned short    rcv_ack_cnt;     /* count of same ack */
    __u32             window_seq;
    __u32             fin_seq;
    __u32             urg_seq;
    __u32             urg_data;
    __u32             syn_seq;
    int               users;           /* user count */
    /*
     * Not all are volatile, but some are, so we
     * might as well say they all are.
     */
    volatile char      dead,
                      urginline,
                      intr,
                      blog,
                      done,
                      reuse,
                      keepopen,
                      linger,
                      delay_acks,
                      destroy,
                      ack_timed,
                      no_check,
                      zapped,
                      broadcast,
                      nonagle,
                      bsdis;
    unsigned long     lingertime;
    int               proc;

    struct sock      *next;
    struct sock      **pprev;
    struct sock      *bind_next;
    struct sock      **bind_pprev;
    struct sock      *pair;
    int               hashent;
    struct sock      *prev;
    struct sk_buff    *volatile send_head;
    struct sk_buff    *volatile send_next;
    struct sk_buff    *volatile send_tail;
    struct sk_buff_head back_log;
    struct sk_buff    *partial;
    struct timer_list partial_timer;
    long              retransmits;
    struct sk_buff_head write_queue,
                      receive_queue;

    struct proto      *prot;
}
```

```

    struct wait_queue      **sleep;
    __u32                  daddr;
    __u32                  saddr;          /* Sending source */
    __u32                  rcv_saddr;      /* Bound address */
    unsigned short         max_unacked;
    unsigned short         window;
    __u32                  lastwin_seq;    /* sequence number when we last
                                           updated the window we offer */
    __u32                  high_seq;      /* sequence number when we did
                                           current fast retransmit */

    volatile unsigned long  ato;           /* ack timeout */
    volatile unsigned long  lrcvtime;     /* jiffies at last data rcv */
    volatile unsigned long  idletime;     /* jiffies at last rcv */
    unsigned int            bytes_rcv;

/*
 *   mss is min(mtu, max_window)
 */
    unsigned short         mtu;           /* mss negotiated in the syn's */
    volatile unsigned short mss;         /* current eff. mss - can change */
    volatile unsigned short user_mss;    /* mss requested by user in ioctl */
    volatile unsigned short max_window;
    unsigned long          window_clamp;
    unsigned int           ssthresh;
    unsigned short         num;
    volatile unsigned short cong_window;
    volatile unsigned short cong_count;
    volatile unsigned short packets_out;
    volatile unsigned short shutdown;
    volatile unsigned long  rtt;
    volatile unsigned long  mdev;
    volatile unsigned long  rto;

    volatile unsigned short backoff;
    int                     err, err_soft; /* Soft holds errors that don't
                                           cause failure but are the cause
                                           of a persistent failure not
                                           just 'timed out' */

    unsigned char          protocol;
    volatile unsigned char  state;
    unsigned char          ack_backlog;
    unsigned char          max_ack_backlog;
    unsigned char          priority;
    unsigned char          debug;
    int                    rcvbuf;
    int                    sndbuf;
    unsigned short         type;
    unsigned char          localroute;    /* Route locally only */

/*
 *   This is where all the private (optional) areas that don't
 *   overlap will eventually live.
 */
    union
    {
        struct unix_opt    af_unix;
#ifdef defined(CONFIG_ATALK) || defined(CONFIG_ATALK_MODULE)
        struct atalk_sock  af_at;
#endif
#ifdef defined(CONFIG_IPX) || defined(CONFIG_IPX_MODULE)
        struct ipx_opt     af_ipx;
#endif
#ifdef CONFIG_INET
        struct inet_packet_opt af_packet;
#endif
#ifdef CONFIG_NUTCP
        struct tcp_opt      af_tcp;
#endif
    } protinfo;

/*
 *   IP 'private area'

```

```

*/
int          ip_ttl;          /* TTL setting */
int          ip_tos;          /* TOS */
struct tcphdr dummy_th;
struct timer_list keepalive_timer; /* TCP keepalive hack */
struct timer_list retransmit_timer; /* TCP retransmit timer */
struct timer_list delack_timer; /* TCP delayed ack timer */
int          ip_xmit_timeout; /* Why the timeout is running */
struct rtable *ip_route_cache; /* Cached output route */
unsigned char ip_hdrincl; /* Include headers ? */
#ifdef CONFIG_IP_MULTICAST
int          ip_mc_ttl; /* Multicasting TTL */
int          ip_mc_loop; /* Loopback */
char         ip_mc_name[MAX_ADDR_LEN]; /* Multicast device name */
struct ip_mc_socklist *ip_mc_list; /* Group array */
#endif

/*
 * This part is used for the timeout functions (timer.c).
 */
int          timeout; /* What are we waiting for? */
struct timer_list timer; /* This is the TIME_WAIT/receive
 * timer when we are doing IP
 */

struct timeval stamp;

/*
 * Identd
 */
struct socket *socket;

/*
 * Callbacks
 */
void          (*state_change)(struct sock *sk);
void          (*data_ready)(struct sock *sk, int bytes);
void          (*write_space)(struct sock *sk);
void          (*error_report)(struct sock *sk);

};

```

## socket

Each socket data structure holds information about a BSD socket. It does not exist independently; it is, instead, part of the VFS inode data structure.

```

struct socket {
    short          type;          /* SOCK_STREAM, ... */
    socket_state   state;
    long          flags;
    struct proto_ops *ops; /* protocols do most everything */
    void          *data; /* protocol data */
    struct socket *conn; /* server socket connected to */
    struct socket *iconn; /* incomplete client conn.s */
    struct socket *next;
    struct wait_queue **wait; /* ptr to place to wait on */
    struct inode *inode;
    struct fasync_struct *fasync_list; /* Asynchronous wake up list */
    struct file *file; /* File back pointer for gc */
};

```

## task\_struct

Each task\_struct data structure describes a process or task in the system.

```

struct task_struct {
    /* these are hardcoded - don't touch */
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */
    long counter;

```

```

    long                priority;
    unsigned            long signal;
    unsigned            long blocked;    /* bitmap of masked signals */
    unsigned            long flags;     /* per process flags, defined below */
    int errno;
    long                debugreg[8];    /* Hardware debugging registers */
    struct exec_domain  *exec_domain;
/* various fields */
    struct linux_binfmt *binfmt;
    struct task_struct *next_task, *prev_task;
    struct task_struct *next_run, *prev_run;
    unsigned long       saved_kernel_stack;
    unsigned long       kernel_stack_page;
    int                 exit_code, exit_signal;
/* ??? */
    unsigned long       personality;
    int                 dumpable:1;
    int                 did_exec:1;
    int                 pid;
    int                 pgrp;
    int                 tty_old_pgrp;
    int                 session;
/* boolean value for session group leader */
    int                 leader;
    int                 groups[NGROUPS];
/*
 * pointers to (original) parent process, youngest child, younger sibling,
 * older sibling, respectively. (p->father can be replaced with
 * p->p_pptr->pid)
 */
    struct task_struct *p_opptr, *p_pptr, *p_cptra,
    *p_ysptr, *p_osptra;
    struct wait_queue  *wait_chldexit;
    unsigned short     uid,euid,suid,fsuid;
    unsigned short     gid,egid,sgid,fsgid;
    unsigned long       timeout, policy, rt_priority;
    unsigned long       it_real_value, it_prof_value, it_virt_value;
    unsigned long       it_real_incr, it_prof_incr, it_virt_incr;
    struct timer_list   real_timer;
    long               utime, stime, cutime, cstime, start_time;
/* mm fault and swap info: this can arguably be seen as either
   mm-specific or thread-specific */
    unsigned long       min_flt, maj_flt, nswap, cmin_flt, cmaj_flt, cnsmap;
    int swappable:1;
    unsigned long       swap_address;
    unsigned long       old_maj_flt;    /* old value of maj_flt */
    unsigned long       dec_flt;        /* page fault count of the last time */
    unsigned long       swap_cnt;       /* number of pages to swap on next pass */
/* limits */
    struct rlimit        rlim[RLIM_NLIMITS];
    unsigned short      used_math;
    char                comm[16];
/* file system info */
    int                 link_count;
    struct tty_struct    *tty;          /* NULL if no tty */
/* ipc stuff */
    struct sem_undo      *semundo;
    struct sem_queue     *semsleeping;
/* ldt for this task - used by Wine. If NULL, default_ldt is used */
    struct desc_struct *ldt;
/* tss for this task */
    struct thread_struct tss;
/* filesystem information */
    struct fs_struct      *fs;
/* open file information */
    struct files_struct  *files;
/* memory management info */
    struct mm_struct      *mm;
/* signal handlers */

```

```

    struct signal_struct *sig;
#ifdef __SMP__
    int                processor;
    int                last_processor;
    int                lock_depth;    /* Lock depth.
                                      We can context switch in and out
                                      of holding a syscall kernel lock... */
#endif
};

```

## timer\_list

timer\_list data structure's are used to implement real time timers for processes.

```

struct timer_list {
    struct timer_list *next;
    struct timer_list *prev;
    unsigned long expires;
    unsigned long data;
    void (*function)(unsigned long);
};

```

## tq\_struct

Each task queue (tq\_struct) data structure holds information about work that has been queued. This is usually a task needed by a device driver but which does not have to be done immediately.

```

struct tq_struct {
    struct tq_struct *next;    /* linked list of active bh's */
    int sync;                 /* must be initialized to zero */
    void (*routine)(void *);  /* function to call */
    void *data;               /* argument to function */
};

```

## vm\_area\_struct

Each vm\_area\_struct data structure describes an area of virtual memory for a process.

```

struct vm_area_struct {
    struct mm_struct * vm_mm;    /* VM area parameters */
    unsigned long vm_start;
    unsigned long vm_end;
    pgprot_t vm_page_prot;
    unsigned short vm_flags;
    /* AVL tree of VM areas per task, sorted by address */
    short vm_avl_height;
    struct vm_area_struct * vm_avl_left;
    struct vm_area_struct * vm_avl_right;
    /* linked list of VM areas per task, sorted by address */
    struct vm_area_struct * vm_next;
    /* for areas with inode, the circular list inode->i_mmap */
    /* for shm areas, the circular list of attaches */
    /* otherwise unused */
    struct vm_area_struct * vm_next_share;
    struct vm_area_struct * vm_prev_share;
    /* more */
    struct vm_operations_struct * vm_ops;
    unsigned long vm_offset;
    struct inode * vm_inode;
    unsigned long vm_pte;        /* shared mem */
};

```

[Top of Chapter](#), [Table of Contents](#), [Show Frames](#), [No Frames](#)

© 1996-1999 David A Rusling [copyright notice](#).