

Classic Problems of Synchronization

- Bounded-Buffer Problem
- Readers-Writers Problem
- Dining Philosophers Problem
- Monitors

Readers-Writers Problem




Readers

Writers


Problem Definition

- Database to be shared among several concurrent processes
- Some processes want to read-only
- Some processes want to read-write
- Several different versions
 - First readers-writers problem
 - Second readers-writers problem

Readers-Writers Problem (Cont.)






Readers-Writers Problem (Cont.)



Concurrent readers

Readers-Writers Problem (Cont.)



First Solution

```
Reader() {  
    while(TRUE) {  
        wait(mutex);  
        readCount++;  
        if(readCount==1)  
            wait(wrt);  
        signal(mutex);  
        read(resource);  
  
        wait(mutex);  
        readCount--;  
        if(readCount == 0)  
            signal(wrt);  
        signal(mutex);  
    }  
}
```

```
Writer() {  
    while(TRUE) {  
        wait(wrt);  
        write(resource);  
        signal(wrt);  
    }  
}  
  
resourceType *resource;  
int readCount = 0;  
semaphore mutex = 1;  
semaphore wrt = 1;
```

- First reader competes with writers
- Last reader signals writers

First Solution (Cont.)

```
Reader() {  
    while(TRUE) {  
        wait(mutex);  
        readCount++;  
        if(readCount == 1)  
            wait(wrt);  
        signal(mutex);  
  
        read(resource);  
  
        wait(mutex);  
        readCount--;  
        if(readCount == 0)  
            signal(wrt);  
        signal(mutex);  
    }  
}
```

```
Writer() {  
    while(TRUE) {  
        wait(wrt);  
        write(resource);  
        signal(wrt);  
    }  
}
```

- First reader competes with writers
- Last reader signals writers
- Any writer must wait for all readers
- Readers can starve writers
- “Updates” can be delayed forever


Second Solution: Writer Precedence

```
Reader() {  
    while(TRUE) {  
        2 → wait(rd);  
            wait(mutex1);  
            readCount++;  
            if(readCount == 1)  
                wait(wrt);  
            signal(mutex1);  
            signal(rd);  
  
            read(resource);  
  
            wait(mutex1);  
            readCount--;  
            if(readCount == 0)  
                signal(wrt);  
            signal(mutex1);  
    }  
}   
int readCount = 0, writeCount = 0;  
semaphore mutex1 = 1, mutex2 = 1;  
semaphore rd = 1, wrt = 1;
```


```
writer() {  
    while(TRUE) {  
        wait(mutex2);  
        writeCount++;  
        1 → if(writeCount == 1)  
            wait(rd);  
        signal(mutex2);  
        3 → wait(wrt);  
  
            write(resource);  
  
        signal(wrt);  
        wait(mutex2);  
        writeCount--;  
        4 → if(writeCount == 0)  
            signal(rd);  
        signal(mutex2);  
    }  
}
```

The Dining Philosophers Problem

- A classical synchronization problem
- 5 philosophers who only eat and think
- Each need to use 2 forks for eating
- There are only 5 forks
- Illustrates the difficulty of allocating resources among process without deadlock and starvation




The Dining Philosophers Problem



- Each philosopher is a process
- One semaphore per fork:
 - Fork: array[0..4] of semaphores
 - Initialization: fork [i].count:=1 for i:=0..4
- A first attempt:
 - Deadlock if each philosopher starts by picking his left fork!

```
Pi () {  
    while (TRUE) {  
        think;  
        wait(fork[i]);  
        wait(fork[i+1 mod 5]);  
        eat;  
        signal(fork[i+1 mod 5]);  
        signal(fork[i]);  
    }  
}
```

The Dining Philosophers Problem



- Idea: admit only 4 philosophers at a time who try to eat
- Then, one philosopher can always eat when the other 3 are holding one fork
- Solution: use another semaphore T to limit at 4 the number of philosophers "sitting at the table"
- Initialize: T.count:=4

```
Pi () {  
    while (TRUE) {  
        think;  
        wait(T) ;  
        wait(fork[i]) ;  
        wait(fork[i+1 mod 5]) ;  
        eat ;  
        signal(fork[i+1 mod 5]) ;  
        signal(fork[i]) ;  
        signal(T) ;  
    }  
}
```


Recall: Problems with Semaphores

- Semaphores are a powerful tool for enforcing mutual exclusion and coordinate processes
- Problem: wait(S) and signal(S) are scattered among several processes
 - It is difficult to understand their effects
 - Usage must be correct in all processes
 - One bad (or malicious) process can fail the entire collection of processes

Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name  
{  
    // shared variable declarations  
    procedure P1 (...) { .... }  
    ...  
    procedure Pn (...) {.....}  
  
    Initialization code ( ....) { ... }  
}
```



Monitors

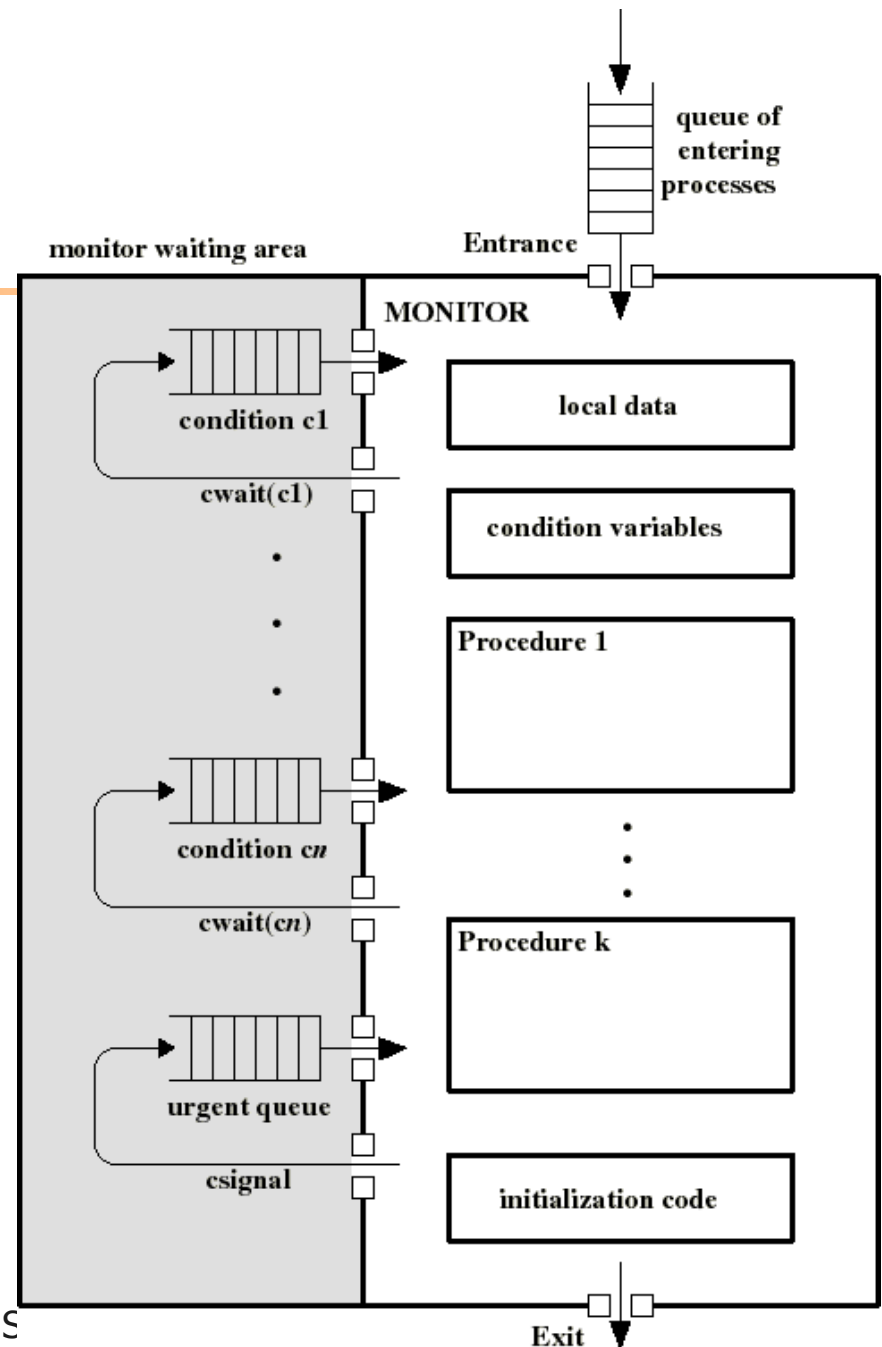
- Is a software module containing:
 - one or more procedures
 - an initialization sequence
 - local shared data variables
- Characteristics:
 - Local shared variables accessible only by monitor's procedures
 - a process enters the monitor by invoking one of its procedures
 - only one process can be in the monitor at any one time
- The monitor ensures mutual exclusion
 - no need to program this constraint explicitly
- Shared data are protected by placing them in the monitor
 - The monitor locks the shared data on process entry

Condition Variables

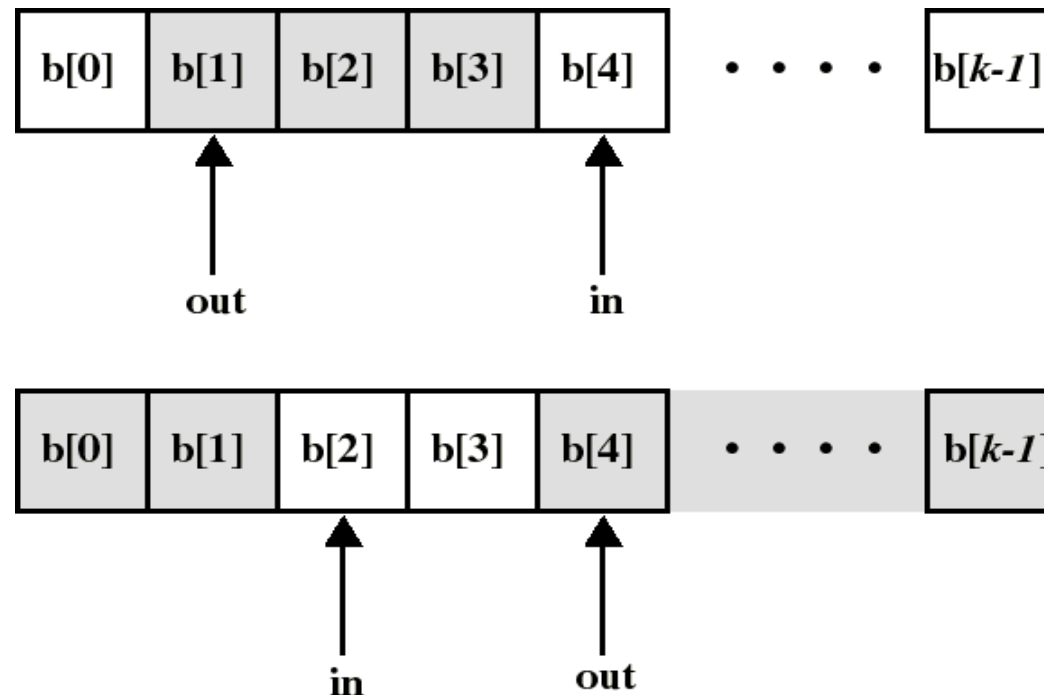
- Process synchronization is done using **condition variables**, which represent *conditions a process may need to wait for before executing in the monitor*
- **condition x, y;**
- Local to the monitor (accessible only within the monitor)
- Can be accessed and changed only by two functions:
 - **x.wait()**: blocks execution of the calling process on condition x
 - the process can resume execution only if another process executes x.signal()
 - **x.signal()**: resume execution of some process blocked on condition x.
 - If several such process exists: choose any one
 - If no such process exists: do nothing

Monitors

- Awaiting processes are either in the entrance queue or in a condition queue
- A process puts itself into condition queue *cn* by issuing *cn.wait()*
- *cn.signal()* brings into the monitor one process in condition *cn* queue
- *signal-and-wait* and *signal-and-continue*



P/C: Finite Circular Buffer of Size k



- Can consume only when the number *full* of (consumable) items is at least 1
- Can produce only when the number *empty* of empty spaces is at least 1

Solution of P/C: Finite Circular Buffer of Size k

Initialization: mutex.count:=1; in:=0;
full.count:=0; out:=0;
empty.count:=k;

Producer:

```
while (TRUE) {  
    produce item;  
    wait(empty);  
    wait(mutex);  
    append(item);  
    signal(mutex);  
    signal(full);  
}
```

Consumer:

```
While (TRUE) {  
    wait(full);  
    wait(mutex);  
    item=take();  
    signal(mutex);  
    signal(empty);  
    consume(item);  
}
```

append(v) :

```
b[in]:=v;  
in:=(in+1)  
    mod k;
```

take() :

```
w:=b[out];  
out:=(out+1)  
    mod k;  
return w;
```

■ critical sections

Producer/Consumer Using Monitors

- Two types of processes:
 - producers
 - consumers
- Synchronization is now confined within the monitor
- `append(.)` and `take(.)` are procedures within the monitor: are the only means by which P/C can access the buffer
- If these procedures are correct, synchronization will be correct for all participating processes

```
Producer:  
while (TRUE) {  
    produce item;  
    append(item) ;  
}
```

```
Consumer:  
while (TRUE) {  
    item=take() ;  
    consume item;  
}
```

Monitor for the Bounded P/C Problem

- Buffer:
 - *buffer*: array[0..k-1] of items;
- Buffer pointers and counts:
 - *nextin*: points to next item to be appended
 - *nextout*: points to next item to be taken
 - *count*: holds the number of items in the buffer
- Condition variables:
 - *notfull*: `notfull.signal()` indicates that the buffer is not full
 - *notempty*: `notempty.signal()` indicates that the buffer is not empty

Monitor for the Bounded P/C Problem

```
Monitor boundedbuffer {
    Item buffer[k];
    integer nextin, nextout, count;
    condition notfull, notempty;

    Append(v) {
        if (count==k) notfull.wait();
        buffer[nextin] = v;
        nextin = (nextin+1) mod k;
        count++;
        notempty.signal();
    }

    initialization_code() {
        nextin=0; nextout=0; count=0;
    }
}

Item Take() {
    if (count==0)
        notempty.wait();
    v = buffer[nextout];
    nextout =
        (nextout+1) mod k;
    count--;
    notfull.signal();
    return v;
}
```

Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses **spinlocks** on multiprocessor systems
- Also provides **dispatcher objects** which may act as either mutexes and semaphores
- Dispatcher objects may also provide **events**
 - An event acts much like a condition variable

Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive
 - Disable kernel preemption on single processors
- Linux provides:
 - semaphores
 - spin locks (mostly on SMP machines)

Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
 - mutex locks
 - condition variables
- Non-portable extensions include:
 - read-write locks
 - spin locks

Transactional Memory

- Multi-core systems
- A **memory transaction**
 - Sequence of memory read-only operations that are atomic
- If all operations are completed, transaction is committed
- Otherwise, rolled back